

iterative optimization

(STOCHASTIC) GRADIENT DESCENT — We seek a hypothesis that is best (among a class \mathcal{H}) according to some notion of how well each hypothesis models given data:

```
def badness(h,y,x):
    # return e.g. whether h misclassifies y,x OR h's surprise at seeing y,x OR etc
def badness_on_dataset(h, examples):
    return np.mean([badness(h,y,x) for y,x in examples])
```

Earlier we found a nearly best candidate by brute-force search over all hypotheses. But this doesn't scale to most interesting cases wherein \mathcal{H} is intractably large. So: *what's a faster algorithm to find a nearly best candidate?*

A common idea is to start arbitrarily with some $h_0 \in \mathcal{H}$ and repeatedly improve to get h_1, h_2, \dots . We eventually stop, say at h_{10000} . The key question is: *how do we compute an improved hypothesis h_{t+1} from our current hypothesis h_t ?*

We *could* just keep randomly nudging h_t until we hit on an improvement; then we define h_{t+1} as that improvement. Though this sometimes works surprisingly well,[◦] we can often save time by exploiting more available information. Specifically, we can inspect h_t 's inadequacies to inform our proposal h_{t+1} . Intuitively, if h_t misclassifies a particular $(x_i, y_i) \in \mathcal{S}$, then we'd like h_{t+1} to be like h_t but nudged toward accurately classifying (x_i, y_i) .[◦]

How do we compute "a nudge toward accurately classifying (x, y) "? That is, how do we measure how slightly changing a parameter affects some result? Answer: derivatives! To make h less bad on an example (y, x) , we'll nudge h in tiny bit along $-g = -\text{dbadness}(h, y, x) / dh$. Say, h becomes $h - 0.01g$.[◦] Once we write

```
def gradient_badness(h,y,x):
    # returns the derivative of badness(h,y,x) with respect to h
def gradient_badness_on_dataset(h, examples):
    return np.mean([gradient_badness(h,y,x) for y,x in examples])
```

we can repeatedly nudge via **gradient descent (GD)**, the engine of ML:[◦]

```
h = initialize()
for t in range(10000):
    h = h - 0.01 * gradient_badness_on_dataset(h, examples)
```

Since the derivative of total badness depends on all the training data, looping 10000 times is expensive. So in practice we estimate the needed derivative based on some *subset* (jargon: **batch**) of the training data — a different subset each pass through the loop — in what's called **stochastic gradient descent (SGD)**:

```
h = initialize()
for t in range(10000):
    batch = select_subset_of(examples)
    h = h - 0.01 * gradient_badness(h, batch)
```

(S)GD requires informative derivatives. Misclassification rate has uninformative derivatives: any tiny change in h won't change the predicted labels. But when we use probabilistic models, small changes in h can lead to small changes in the predicted *distribution* over labels. To speak poetically: the softness of probabilistic models paves a smooth ramp over the intractably black-and-white cliffs of 'right' or 'wrong'. We now apply SGD to maximizing probabilities.

By the end of this section, you'll be able to

- implement gradient descent for any given loss function and (usually) thereby automatically and efficiently find nearly-optimal linear hypotheses from data
- explain why the gradient-update formulas for common linear models are sensible, not just formally but also intuitively

← Also important are the questions of where to start and when to stop. But have patience! We'll discuss these later.

← If you're curious, search 'metropolis hastings' and 'probabilistic programming'.

← In doing better on the i th datapoint, we might mess up how we do on the other datapoints! We'll consider this in due time.

← E.g. if each h is a vector and we've chosen $\text{badness}(h, y, x) = -y h \cdot x$ as our notion of badness, then $-\text{dbadness}(h, y, x) / dh = +yx$, so we'll nudge h in the direction of $+yx$.
Food For Thought: Is this update familiar?

← **Food For Thought:** Can GD directly minimize misclassification rate?

MAXIMUM LIKELIHOOD ESTIMATION — When we can compute each hypothesis h 's asserted probability that the training y s match the training x s, it seems reasonable to seek an h for which this probability is maximal. This method is **maximum likelihood estimation (MLE)**. It's convenient for the overall goodness to be a sum (or average) over each training example. But independent chances multiply rather than add: rolling snake-eyes has chance $1/6 \cdot 1/6$, not $1/6 + 1/6$. So we prefer to think about maximizing log-probabilities instead of maximizing probabilities — it's the same in the end.[◦] By historical convention we like to minimize badness rather than maximize goodness, so we'll use SGD to *minimize negative-log-probabilities*.

← Throughout this course we make a crucial assumption that our training examples are independent from each other.

```
def badness(h,y,x):
    return -np.log( probability_model(y,x,h) )
```

Let's see this in action for the linear logistic model we developed for soft binary classification. A hypothesis \vec{w} predicts that a (featurized) input \vec{x} has label $y = +1$ or $y = -1$ with chance $\sigma(+\vec{w} \cdot \vec{x})$ or $\sigma(-\vec{w} \cdot \vec{x})$:

$$p_{y|x,w}(y|\vec{x}, \vec{w}) = \sigma(y\vec{w} \cdot \vec{x}) \quad \text{where} \quad \sigma(\vartheta) = 1/(1 - \exp(-\vartheta))$$

So MLE with our logistic model means finding \vec{w} that *minimizes*

$$-\log(\text{prob of all } y_i\text{'s given all } \vec{x}_i\text{'s and } \vec{w}) = \sum_i -\log(\sigma(y_i\vec{w} \cdot \vec{x}_i))$$

The key computation is the derivative of those badness terms:[◦]

← Remember that $\sigma'(z) = \sigma(z)\sigma(-z)$. To reduce clutter we'll temporarily write $y\vec{w} \cdot \vec{x}$ as ywx .

$$\frac{\partial(-\log(\sigma(ywx)))}{\partial w} = \frac{-\sigma(ywx)\sigma(-ywx)yx}{\sigma(ywx)} = -\sigma(-ywx)yx$$

Food For Thought: If you're like me, you might've zoned out by now. But this stuff is important, especially for deep learning! So please graph the above expressions to convince yourself that our formula for derivative makes sense visually.

To summarize, we've found the loss gradient for the logistic model:

```
sigma = lambda z : 1./(1+np.exp(-z))
def badness(w,y,x):      return -np.log( sigma(y*w.dot(x)) )
def gradient_badness(w,y,x): return -sigma(-y*w.dot(x)) * y*x
```

As before, we define overall badness on a dataset as an average badness over examples; and for simplicity, let's initialize gradient descent at $h_0 = 0$:

```
def gradient_badness_on_dataset(h, examples):
    return np.mean([gradient_badness(h,y,x) for y,x in examples])
def initialize():
    return np.zeros(NUMBER_OF_DIMENSIONS, dtype=np.float32)
```

Then we can finally write gradient descent:

```
h = initialize()
for t in range(10000):
    h = h - 0.01 * gradient_badness_on_data(h, examples)
```